Measured training throughput on one and two H100 GPUs

Same microbatch: dual / single = 1.92× Best dual / single = 3.14× (also changes microbatch)

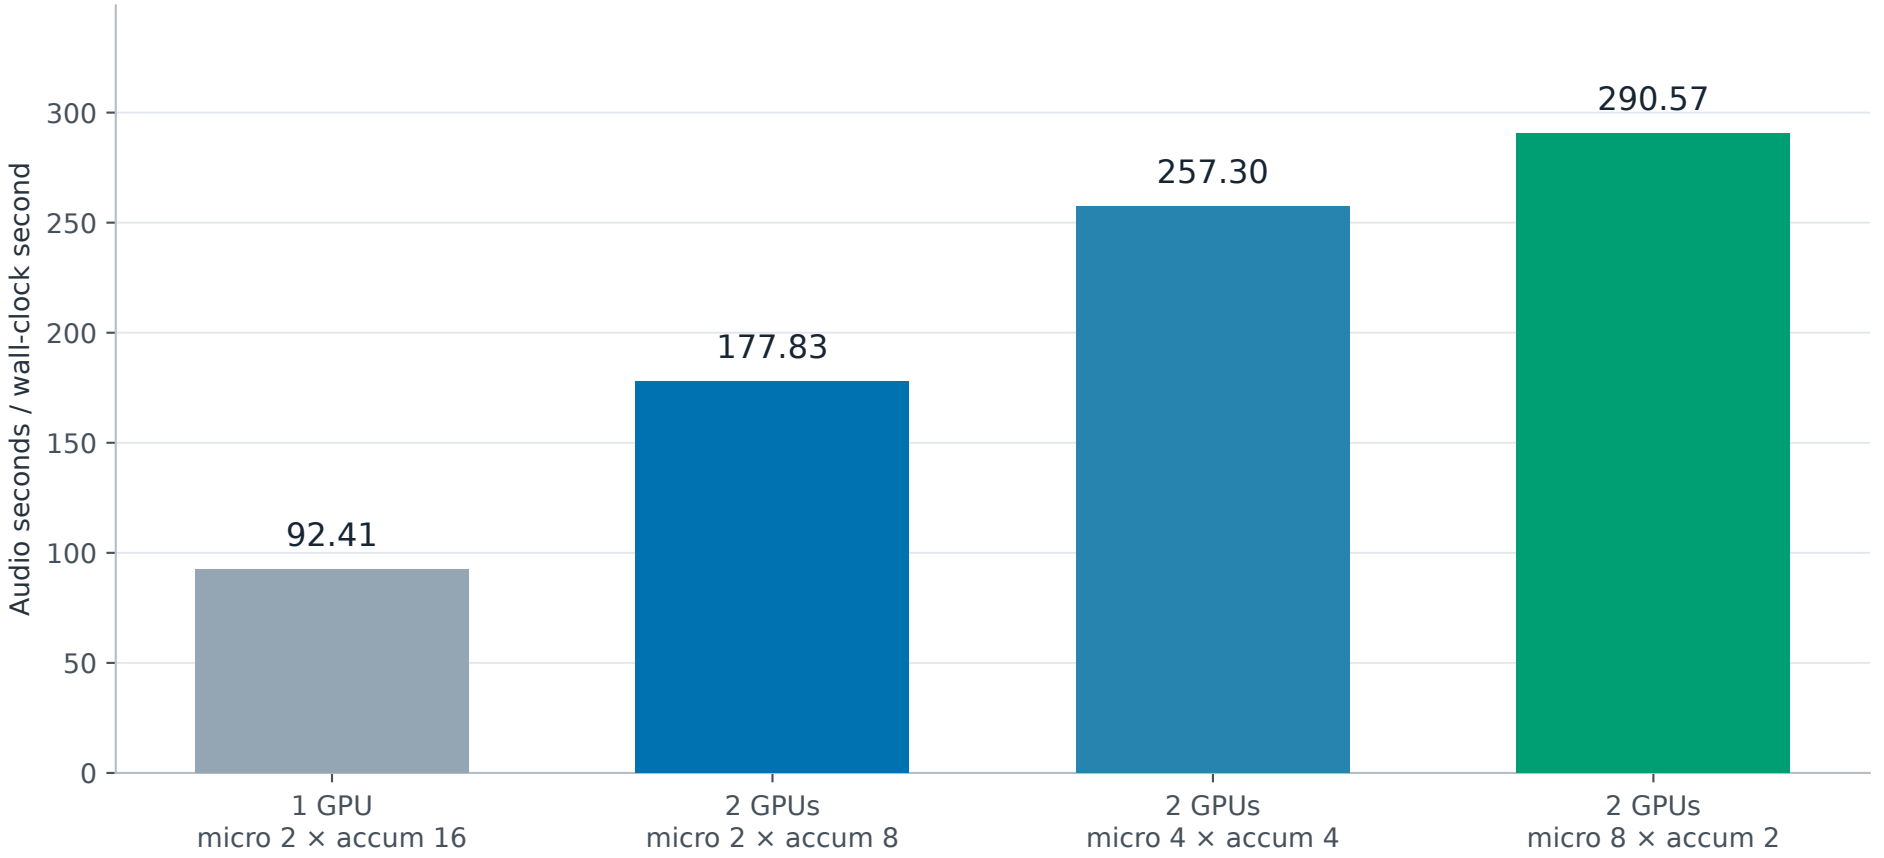

Same pinned model, 2,048-example subset and global batch 32; updates 11–40 after 10 warmup updates.
Update timing excludes model loading, checkpoint writes and generation evaluation. GPU0 also hosts PersonaPlex; NODE topology, no assumed NVLink.