Qwen3-ASR: observed training dynamics

Full-parameter SFT; original run: global batch 32, learning rate 1e-5, one epoch

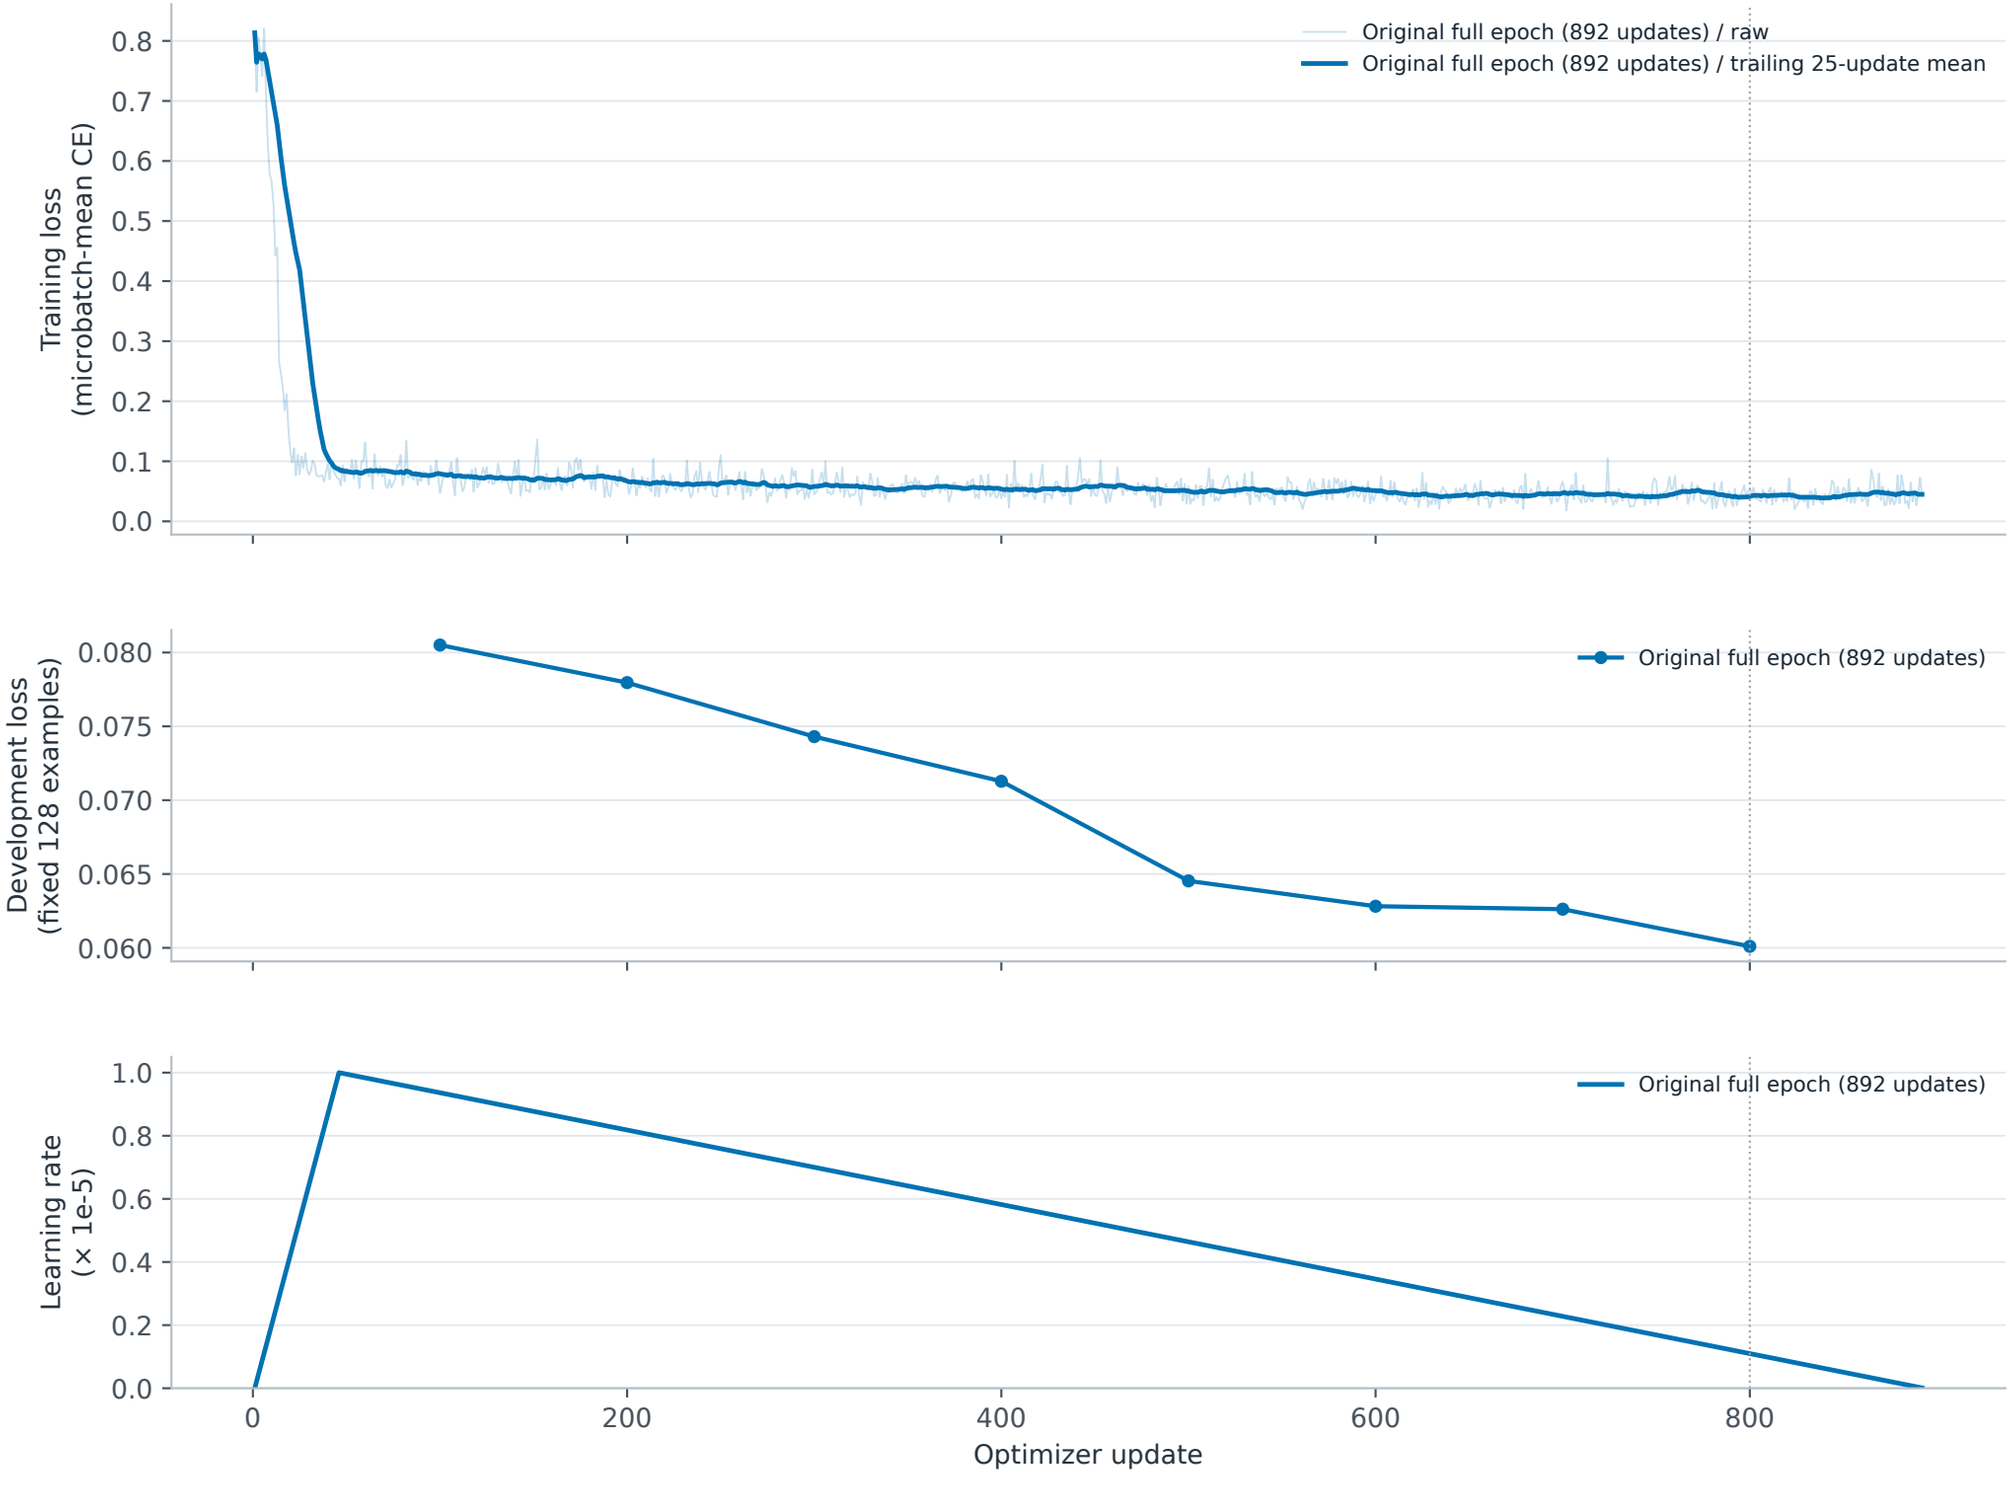

Original run: 892 training records; its development loss is observed only at updates 100–800. No endpoint dev-loss point is invented.
Loss is not generation WER. Raw losses and explicitly labeled trailing means are both shown; learning-rate values are logged observations.